

# IDS Projects

Vania Marangozova-Martin

2020-2021

# Schedule


S12	22/3		7	Discuss LABS, Choice of project, P2P		1.5	////		
S13	29/03		8	P2P II, Complex dist systems I		3	////		
S14	5/4		9				Project - design 1st version		1.5
S15	12/4		10	Complex dist systems II		1.5	Project		3
S16	19/4		Spring VACATION   ** TURN IN PROJECT **				** TURN IN PROJECT **		
S17	26/4		11	SOUTENANCES		1.5	SOUTENANCES		6
S18	3/5		EXAMENS				15.		18.

# Administrativa

- Organisation
  - 2 members per group
- Project advancement:
  - We will organize discussions per group
- Turn in your project
  - turn in code and a report
  - prepare live remote demo session (15 minutes)
  - report = description of architecture, design, choice of techno
- There are three "algo-oriented projects"
  1. you need to work on the algorithmic logic
  2. implement it with the techno (RMI or RabbitMQ) you want
  3. be capable of explaining the choices (algo/techno) you made

# Project 1 [ALGO]: Network Overlay

1. The goal is to construct a virtual network over a physical one
2. We start with a graph that gives the physical topology and build a virtual topology




# Overlay : the input graph

1. You need an input to describe a graph for example a matrix
2. As the graph gives the physical topology of your system, you need to instantiate it with RabbitMQ or Java RMI to have a running distributed system

	1	2	3	4	5
1			1		
2			1	1	1
3	1	1			
4		1			
5		1			



input matrix




running system, each node is a separate running entity (e.g. a remote object in RMI)

# Overlay : virtual ring

- Decide of the nodes identifiers




different choice





# Overlay : virtual ring

- Decide who is neighbor of who in the virtual ring
  - May be 1-2-3-4-5-1 or, as given in the picture, 1-4-3-5-2-1



# Overlay : virtual ring



- Compute the routing between a ring node and its two neighbors





# Overlay : virtual ring

- Implement the two basic primitives  
`sendLeft(Message m)` and `sendRight(Message m)`



# Project 2 [ALGO]: Multi-player Game

- Goal : Follow the movements of players
- Setting
  - The playground is separated in zones
    - you need to decide how
  - Players move around and pass from one zone to another
    - how do they move?
    - should avoid collisions
    - should say "Hello" to each other
  - Each zone is managed by a node
  - The nodes that manage the set of zones are connected in a distributed topology
    - You need to choose what topology

# Project 2 [ALGO]: Multi-player Game (2)


## EXAMPLE of a setting

playground




# Project 2 [ALGO]: Multi-player Game (2)

## EXAMPLE of a setting




# Project 2 [ALGO]: Multi-player Game (2)

## EXAMPLE of a setting




# Project 2 [ALGO]: Multi-player Game (2)

## EXAMPLE of a setting




## Project 2 [ALGO]: Multi-player Game (2)

EXAMPLE: The corresponding dist. architecture



# Project 2 : Multi-player Game (2)


Red says Hello to Orange





# Project 2 : Multi-player Game (2)

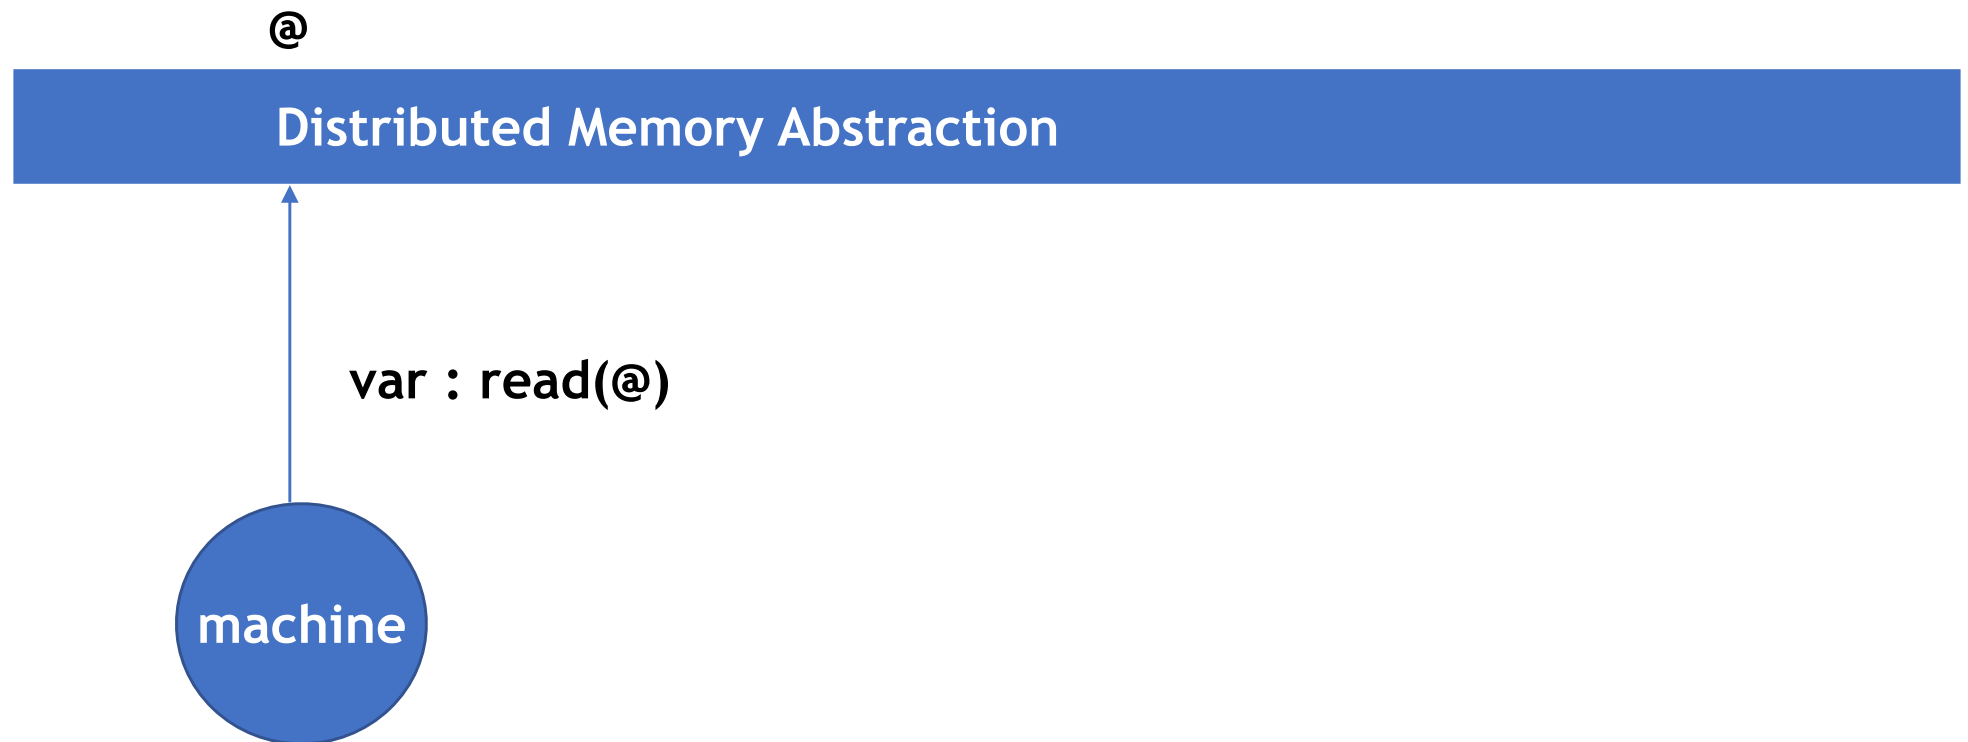
Red says Hello to Orange



Green enters zone C

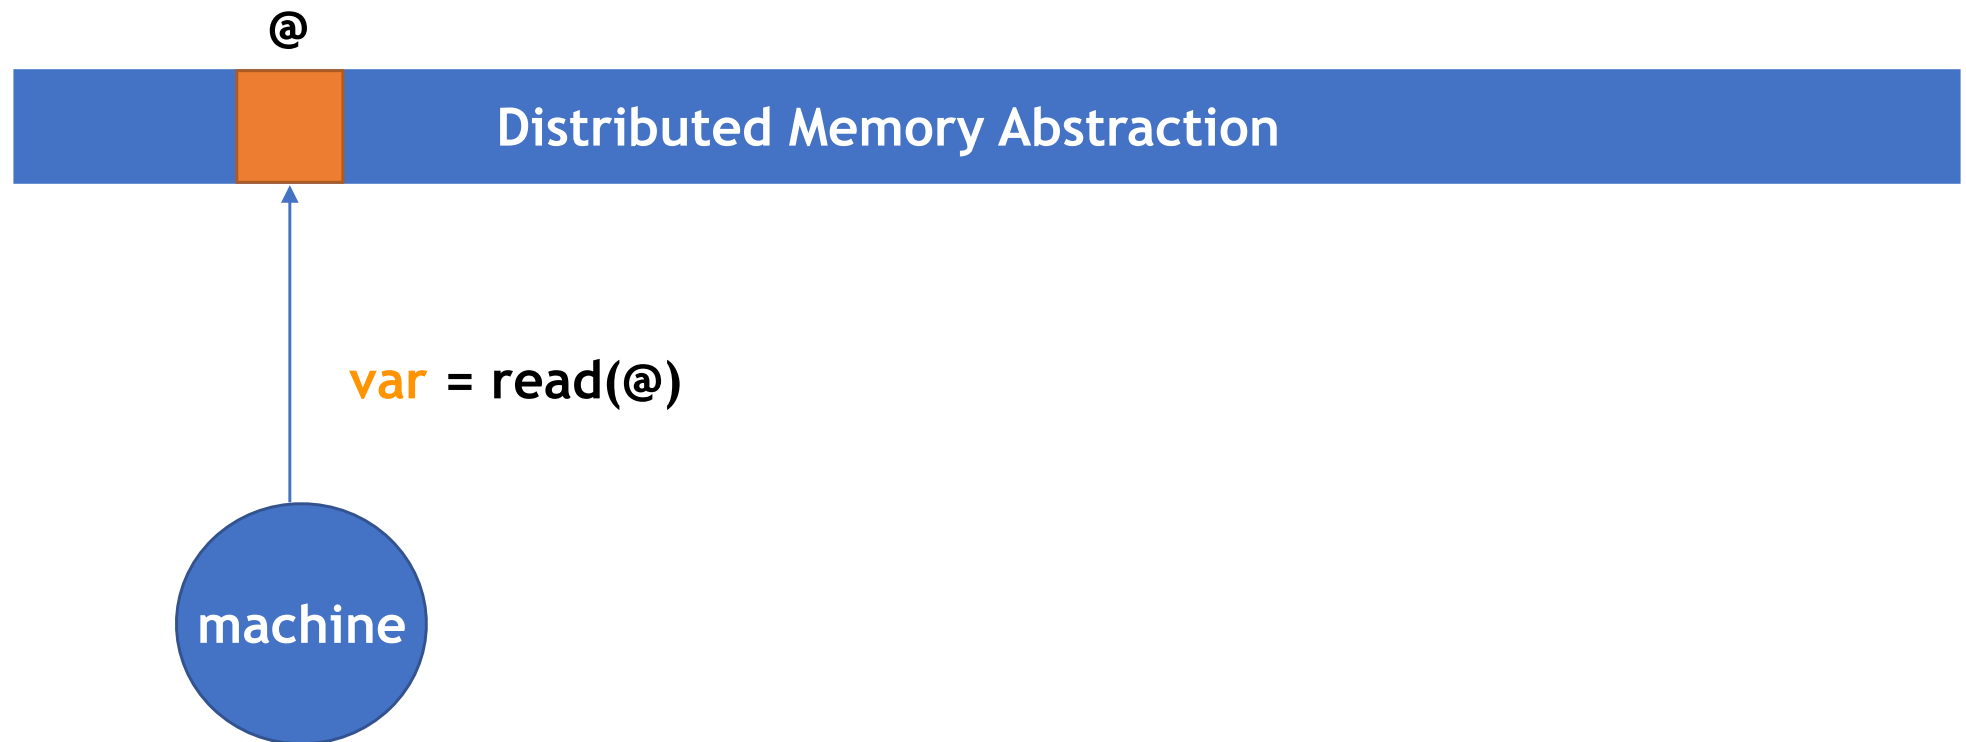
# Project 3 [ALGO]: DSM (Distributed Shared Memory)

- Distributed Shared Memory Abstraction



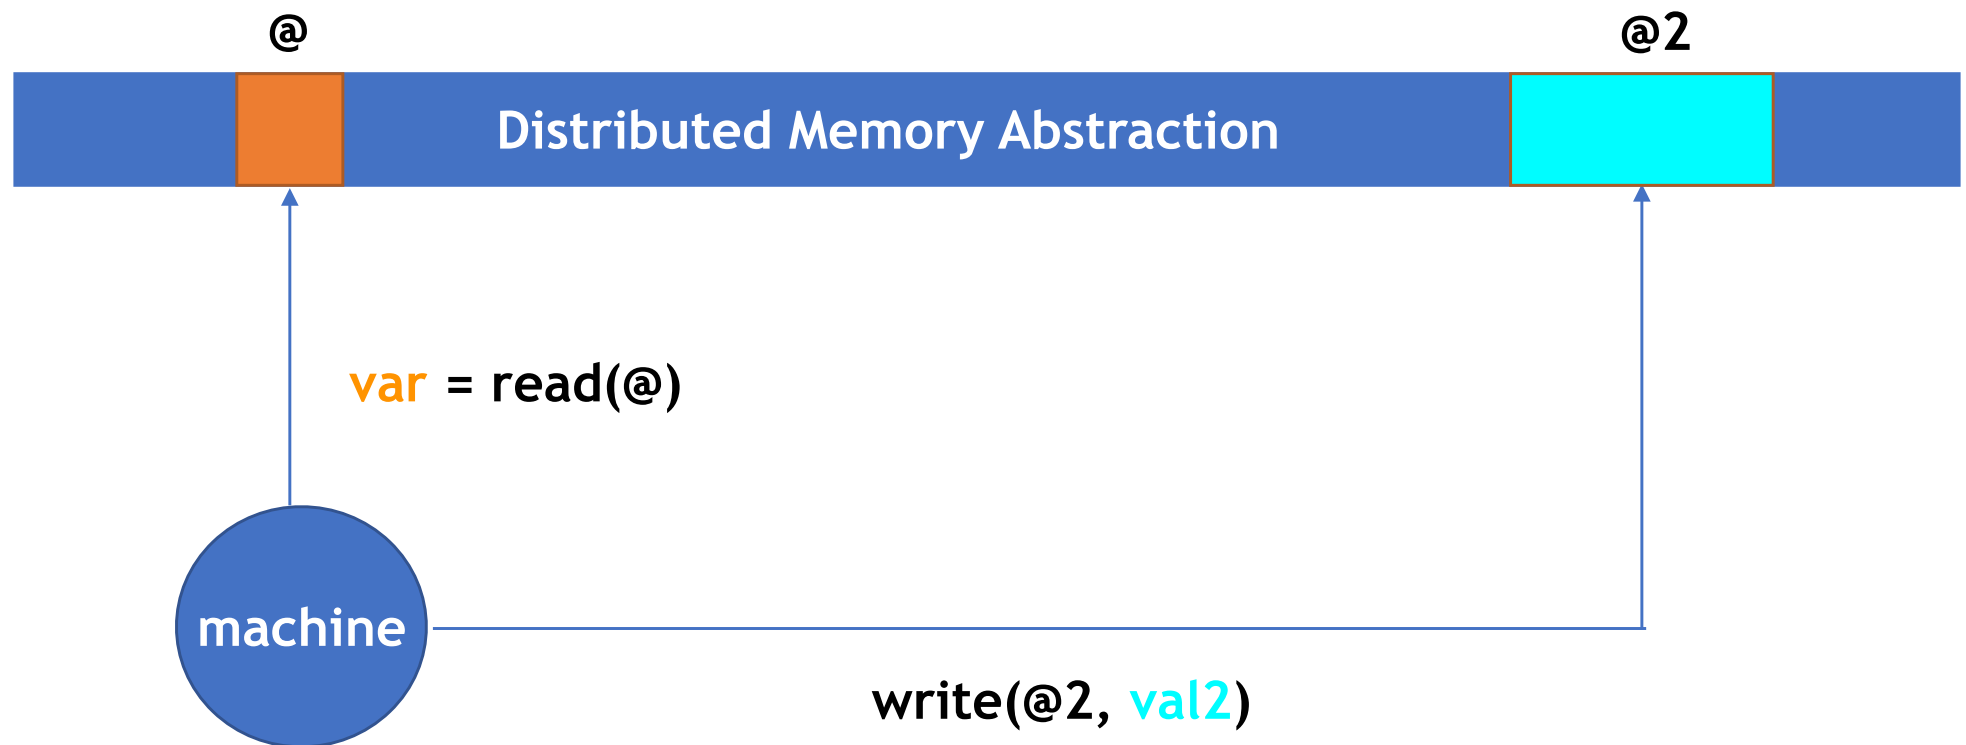
# Project 3 [ALGO]: DSM (Distributed Shared Memory)

- Distributed Memory Abstraction



# Project 3 [ALGO]: DSM (Distributed Shared Memory)

- Distributed Memory Abstraction



# Project 3 [ALGO]: DSM (Distributed Shared Memory)

To Implement the primitives

- `value read(address)`
- `write(address, value)`


1. You need to decide **what is an address** and how to manage addresses
2. You need to decide **what kind of values** you will manipulate

# Project 3 [ALGO]: DSM (Distributed Shared Memory)

- The memory is actually distributed over several physical machines
  - You need to decide how to distribute the memory
- The machines are interconnected into a distributed topology
  - You need to choose the topology
- You need to decide what is an address and how to manage addresses
- You need to decide what kind of values you will manipulate


# Project 3 [ALGO]: DSM EXAMPLE Setting

## Distributed Shared Memory Abstraction




**Physical machines in a grid**

# Project 3 [ALGO]: DSM EXAMPLE with Partitioning






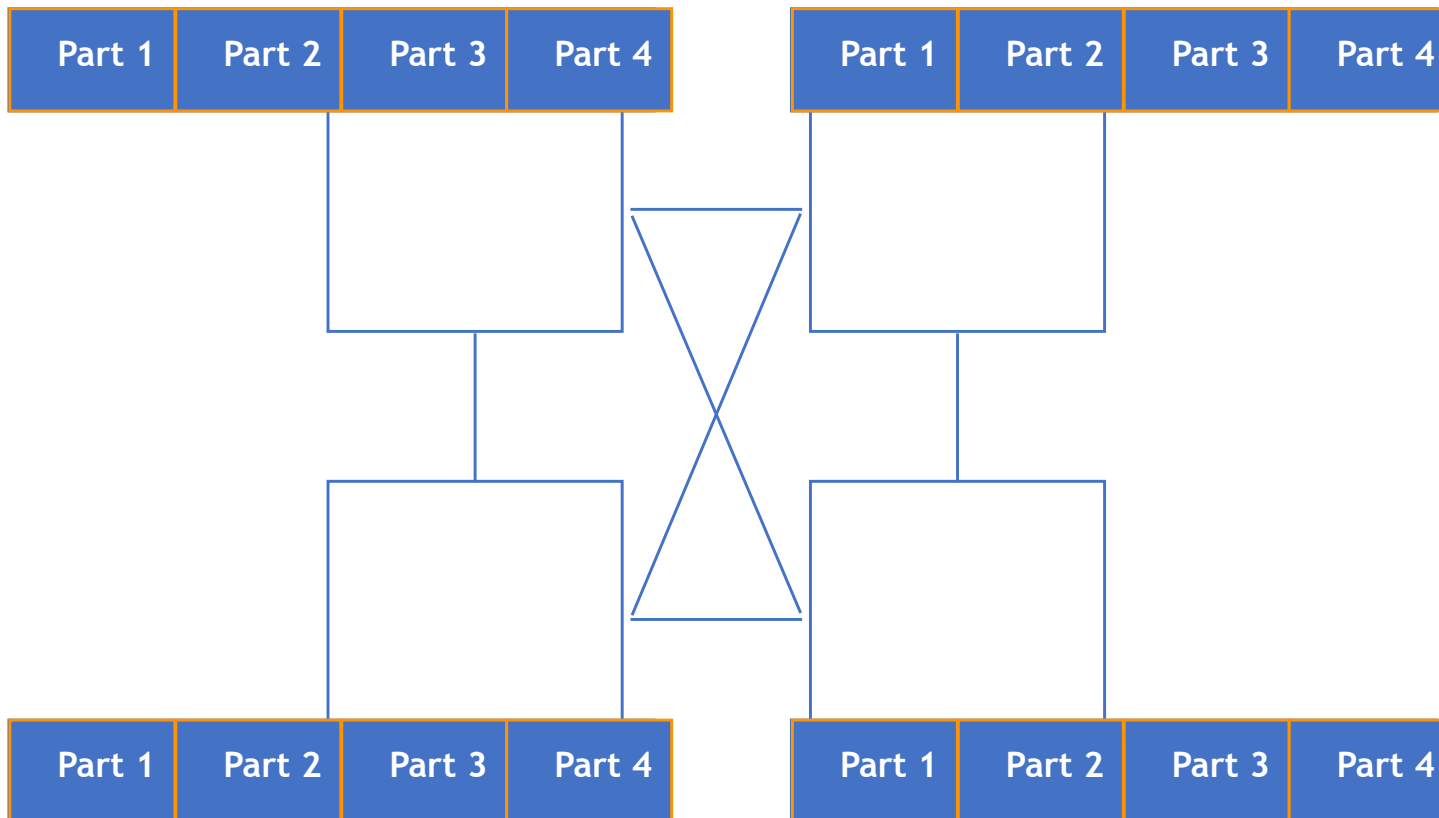
# Project 3 [ALGO]: DSM EXAMPLE with Partitioning



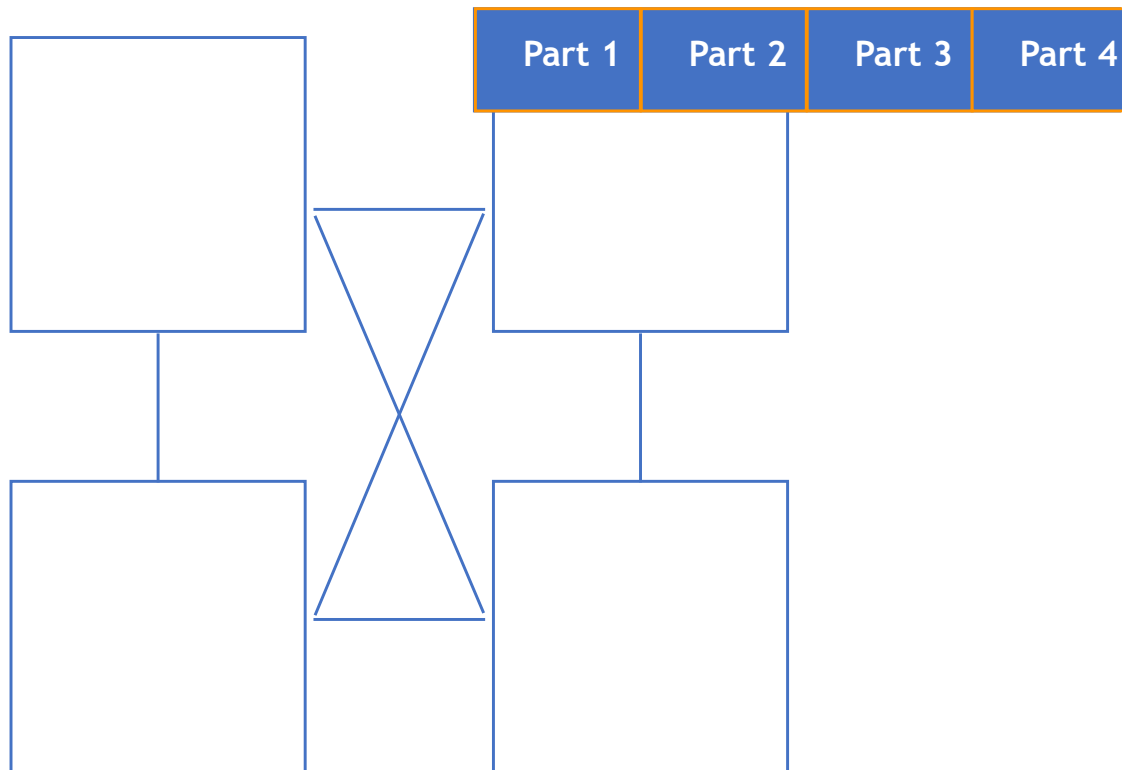
# Project 3 [ALGO]: DSM EXAMPLE with Partitioning



# Project 3 [ALGO]: DSM EXAMPLE with Replication



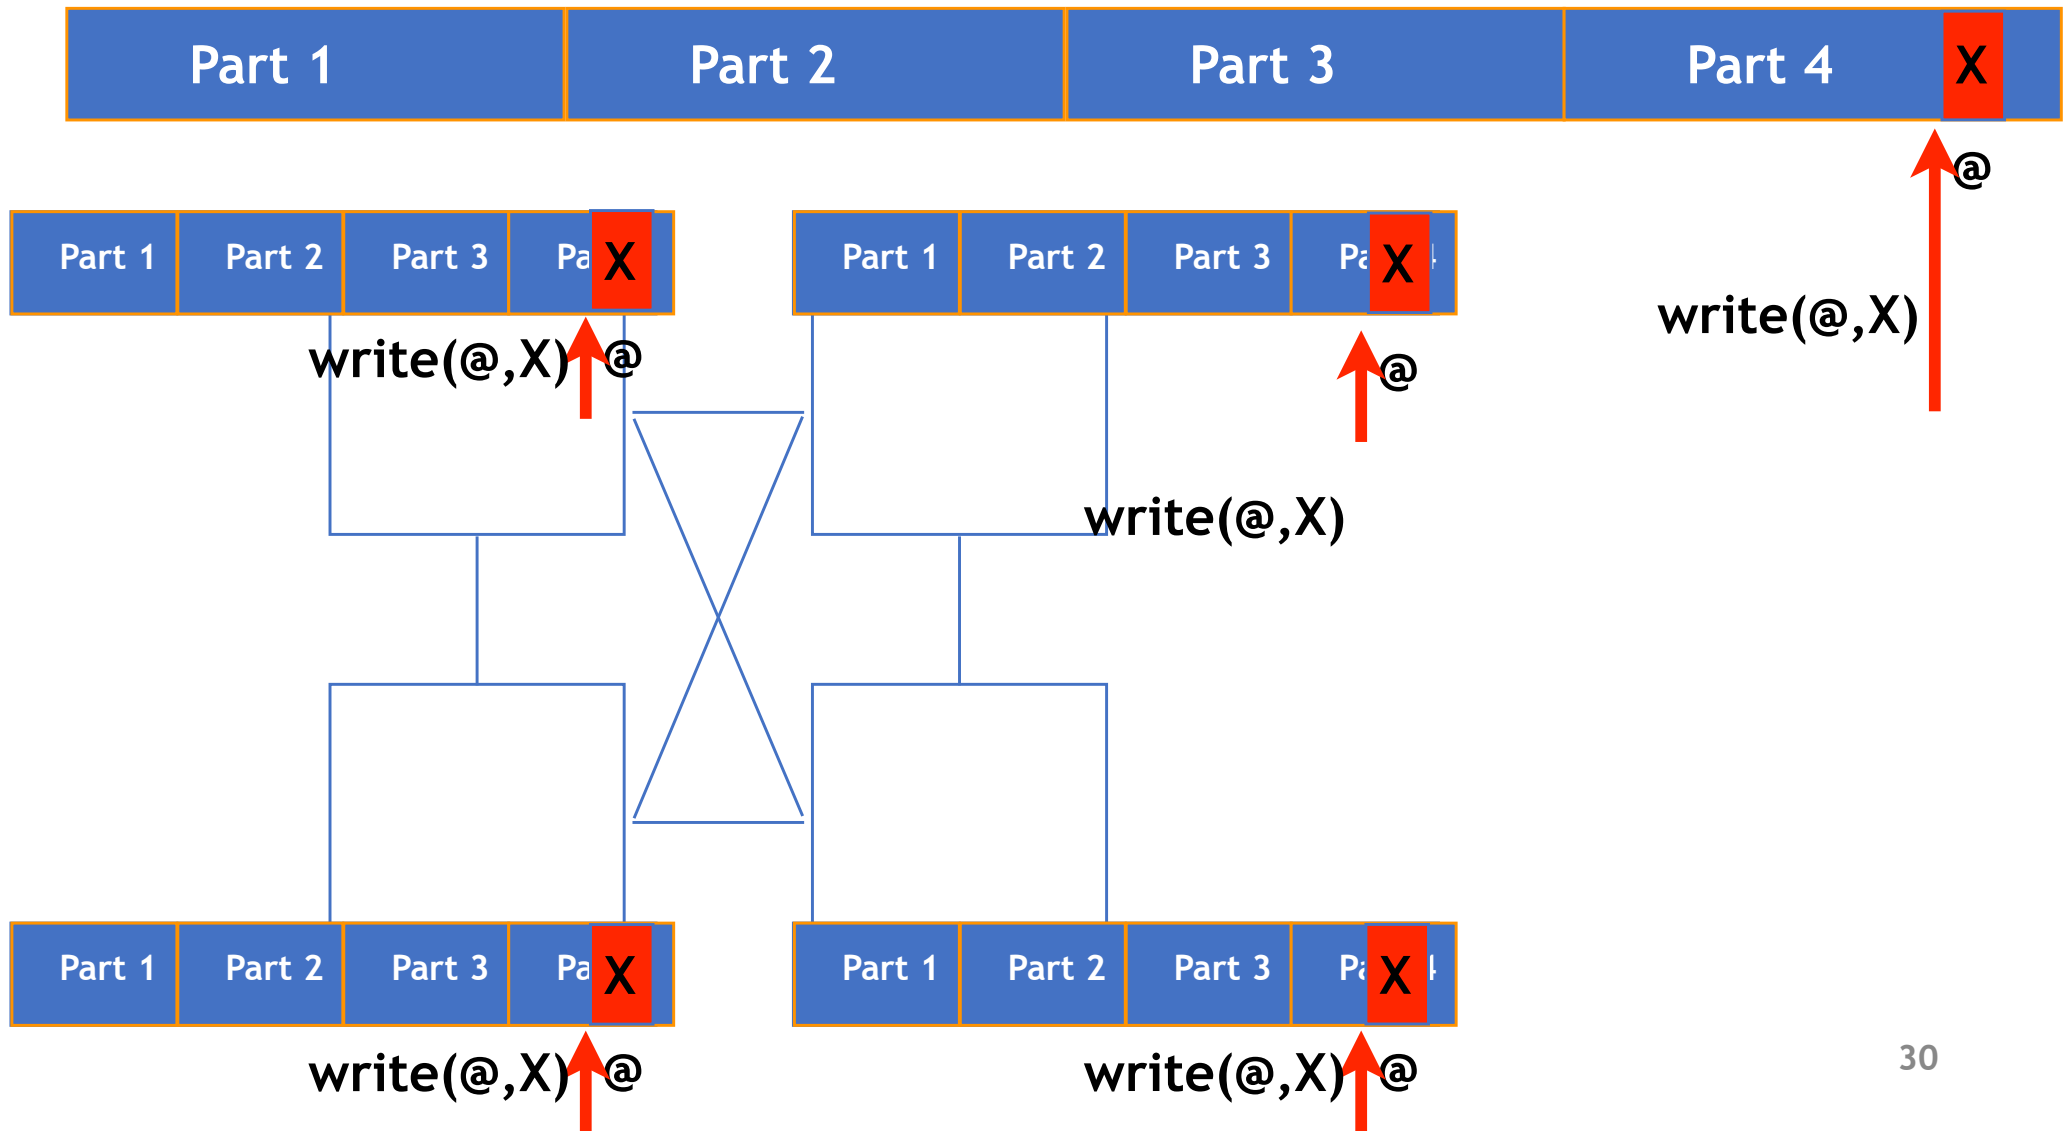
# Project 3 [ALGO]: DSM EXAMPLE with Non Replication



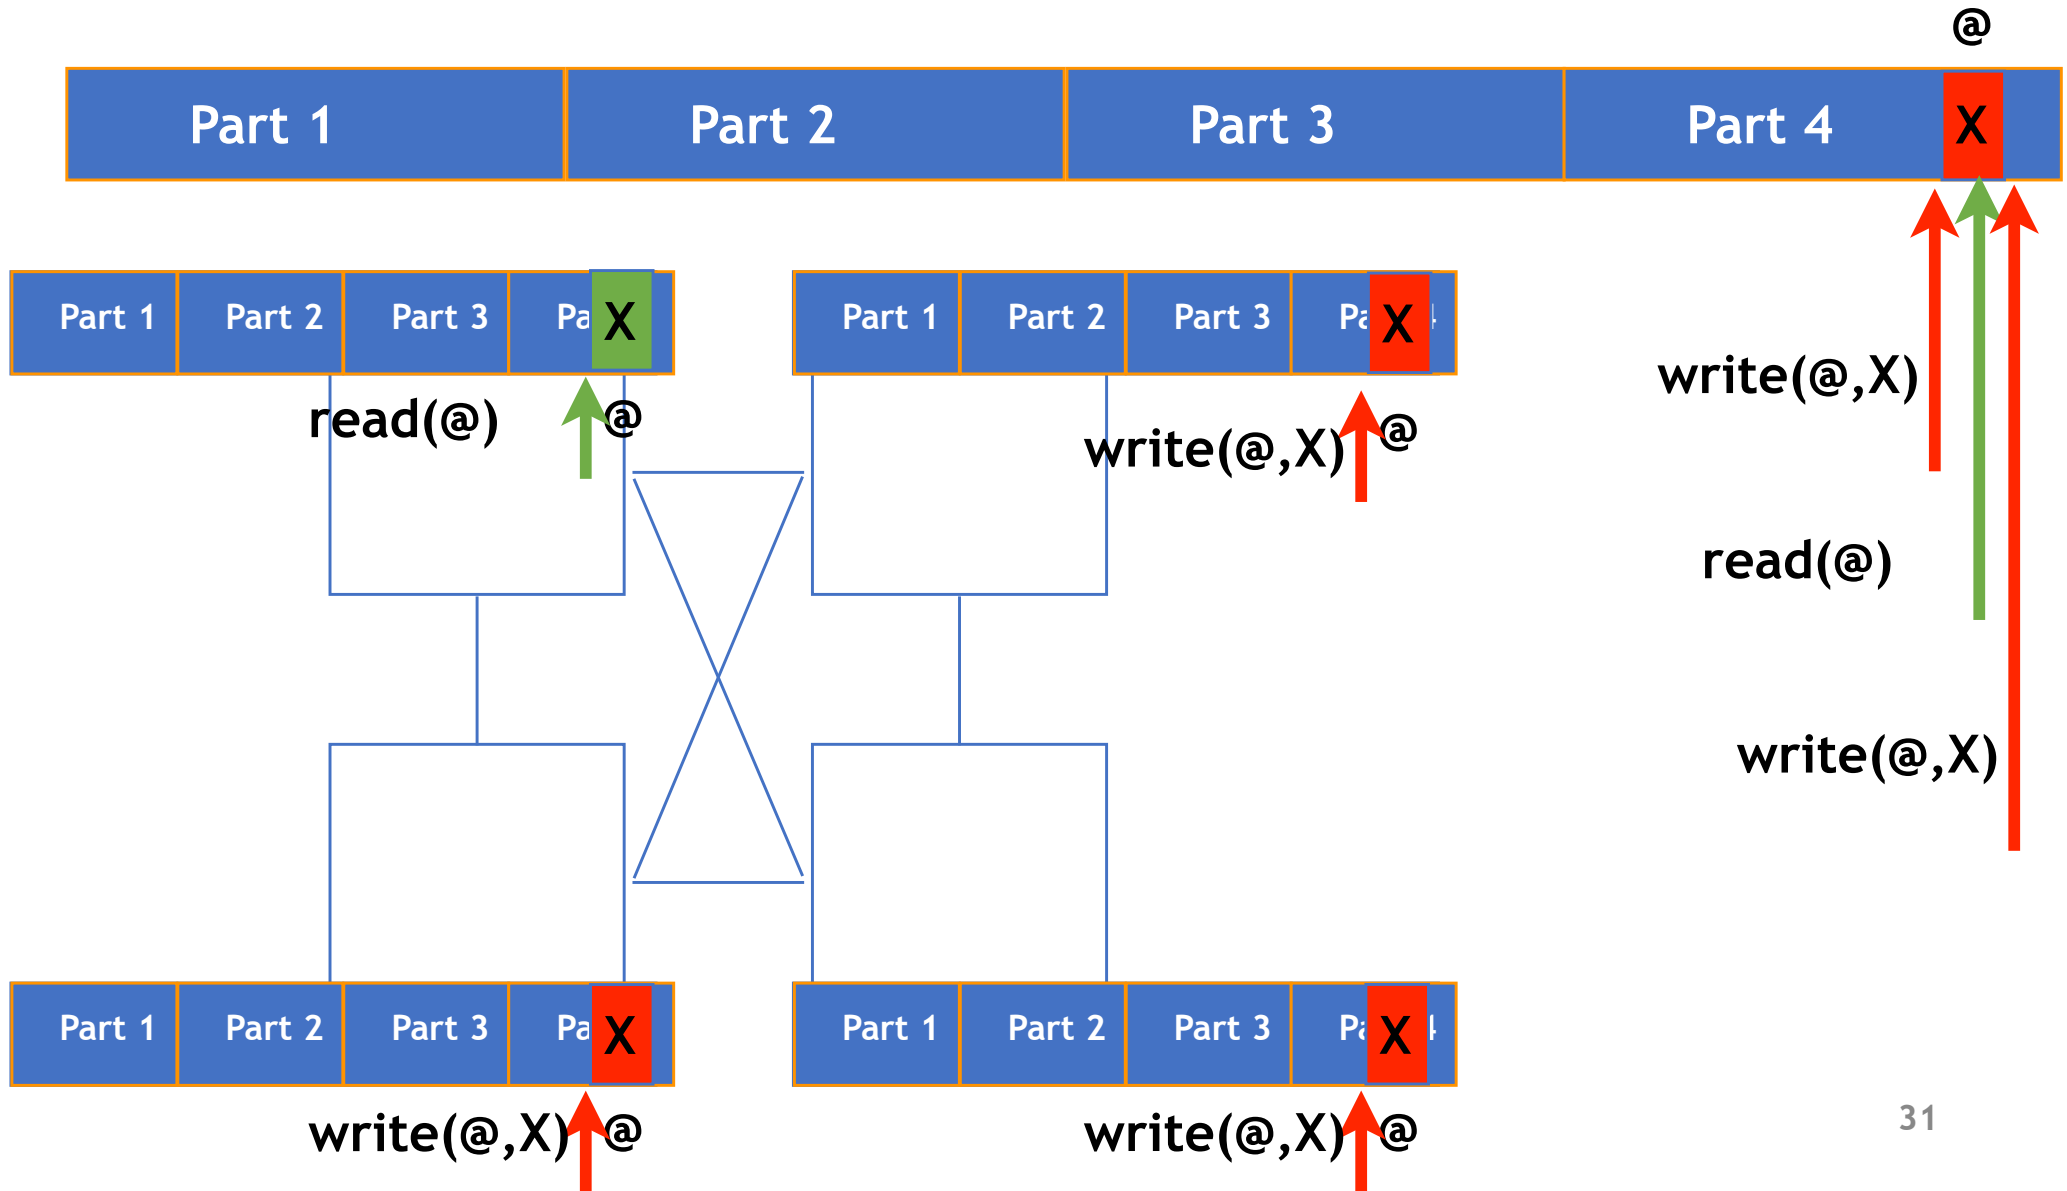
# The big problem is **consistency** (this is really complicated)

- If there are only read operations, no problem
- The pb is that there are write operations
- What if two write operations concern the same address?
  - Which one will go first?
  - What value will the others observe?  
(if they read from this address)

# Strong consistency



# Consistency?



# References (on Moodle)

1. Fekete, A. D., Ramamritham, K.: Consistency Models for Replicated Data. In Charron- Bost, B., Pedone, F., Schiper, A., eds. : Replication LNCS 5959. Springer-Verlag, Berlin (2010) 1-17
2. Burckhardt, S.: Principles of Eventual Consistency. Foundations and Trends Programming Languages 1(1-2), 1-150 (2014)